



STUDENT PARTICIPATION DATASET ANALYSIS

Week 4 - Final Program Insights Report

Prepared by Team 8.

Internship program: AI Data Internship

1. Overview

1.1 Project Background

This Week 4 report is the final deliverable of the RIT Opportunity-Wise Data Analysis project — a four-week structured analytics internship conducted under the Excelerate AI Data Internship programme. The project analysed 8,558 student participation records across 22 opportunities on the RIT platform, progressing from raw data documentation through Exploratory Data Analysis, machine learning modelling, and finally to this synthesised programme insights and recommendations report.

1.2 Four-Week Project Arc

Week	Focus	Key Deliverable
Week 1	Data Understanding & Documentation	Data dictionary, dataset overview, initial cleaning steps, summary statistics
Week 2	Exploratory Data Analysis (EDA)	Participation trends, demographic patterns, temporal analysis, anomaly identification
Week 3	Machine Learning & Predictive Modelling	3 trained models, probability scores for 22 opportunities, feature importance, heatmaps
Week 4	Final Insights & Recommendations	This report: synthesis, prioritised recommendations, measurement framework, reflection

1.3 Tools & Technologies

Python (pandas, scikit-learn, matplotlib, seaborn) · Microsoft Word · PowerPoint · Google Drive

2. Executive Summary

Across four weeks of rigorous analysis, Team 8 investigated the participation patterns of 3,026 learners applying to 22 opportunities across 5 categories on the RIT platform. The central finding is both simple and actionable: the platform has a supply problem, not a learner quality problem. Internship-type opportunities attract 63.3% of all applications but deliver only a 21.5% Active participation rate, not because learners are unsuitable, but because placement capacity is insufficient.

Three machine learning models were trained to predict Active vs Inactive participation, all achieving accuracy above 86% and AUC-ROC above 90%. Gradient Boosting (AUC 91.99%) is recommended for deployment. The dominant predictor, Opportunity Category at 52–55% feature importance, confirms the structural nature of the participation gap. Predicted probabilities for all 22 opportunities were generated, ranging from 99.6% (Engagement formats) to 29.9% (Data Visualization Internship).

Six prioritised recommendations are presented, covering internship capacity expansion, outreach timing, Africa and gender equity, drop-out prevention, course pipeline design, and a four-phase ML deployment roadmap. A measurement framework with specific KPI targets is provided to enable Excelerate to track programme improvement across future cohorts.

3. Summary of Findings

3.1 Participation Trends

The dataset contains 8,558 application records, of which 47.2% are Active, 49.9% are Inactive, 2.5% are Pending, and 0.3% are Completed. The rejection rate across all applications is 41.7%, with Team Allocated representing 38.3%, fewer than four in ten applicants are successfully placed. Internship drop-out (617 records, 7.2% of all applications) is concentrated almost entirely within the Internship category (613 of 617, 99.4%).

Metric	Value	Significance
Total Records	8,558	Full dataset including all application stages
Active Rate (overall)	47.2%	Less than half of all applications result in active participation
Rejection Rate	41.7%	Primary inactive outcome — driven by internship capacity constraints
Internship Share of Apps.	63.3%	Highest demand category with lowest Active rate (21.5%)
Non-internship Active Rate	92–97%	Courses, Events, Competitions, Engagement retain nearly all participants
Drop-out Concentration	99.4% in Internship	Structural attrition risk concentrated in one category
Q1 2024 Surge	2,829 apps (Jan–Feb)	10× monthly baseline — evidence of outreach effectiveness

3.2 Demographic Insights

- Age:** 64.5% of learners are aged 21–25 (undergraduate/early postgraduate). Learners 16–20 show lower internship predicted Active probability (0.38 vs 0.44 for 21–25).

- **Gender:** 58.6% Male, 41.2% Female. Gender is the lowest-importance predictor globally (2.5%) but shows amplified gaps in Africa, female learners there have notably lower Active rates than male learners.
- **Region:** North America (46.5%) and Asia (37.4%) dominate. Africa (14.8%) is the third-largest and fastest-growing region, with Nigeria alone contributing 8.9% of all records.
- **STEM dominance:** Information Systems (2,158 records) and Computer Science (1,052) are the most represented majors, together accounting for 37% of all applications.

3.3 Predictive Model Results

Three supervised machine learning classifiers were trained on 8,315 records (Active vs Inactive, 80/20 stratified split) using six features: Opportunity Category, Gender, Region, Age Group, Application Quarter, and Learner Age.

Model	Accuracy	Precision	Recall	F1-Score	AUC-ROC	CV AUC (5-Fold)
Random Forest	86.29%	85.68%	86.29%	85.88%	91.70%	91.46% $\pm$ 0.85%
Gradient Boosting ✓	86.71%	86.26%	86.40%	86.27%	91.99%	91.12% $\pm$ 0.87%
Logistic Regression	86.05%	85.67%	85.67%	85.67%	90.89%	89.58% $\pm$ 1.16%

Gradient Boosting is the recommended model. It leads on every metric, with the most stable cross-validation performance ($\pm 0.87\%$), confirming reliable generalisation to new applicant data.

3.4 Feature Importance




Figure 1: Feature importance — Random Forest (left) vs Gradient Boosting (right). Opportunity Category dominates at 52–55%.

Opportunity Category accounts for 52–55% of all predictive information. Region (12–14%) and Application Quarter (11–13%) are the next most informative. Gender contributes only 2.5% globally but shows local amplification in African sub-groups. These results confirm that participation is predominantly determined by the type of opportunity offered, not by the demographics of the applicant.

3.5 Predicted Participation Probabilities




Figure 2: Ensemble-average predicted Active participation probability for all 22 opportunities (RF + GB average).

Non-internship opportunities (Engagement, Event, Course, Competition) predict 96–99.6% Active probability, near-certain participation. All internship-type opportunities fall below 60%, with Data Visualization Internship at 29.9% representing the most severe capacity gap.

3.6 Predicted Probability Heatmaps




Figure 3: Predicted Active Probability — Opportunity Category × Region (Random Forest).




Figure 4: Predicted Active Probability — Opportunity Category × Age Group (Random Forest).

The Category × Region heatmap confirms that non-internship Active rates are consistent across all geographies (0.97–0.99). For Internships, North America (0.42) and Asia (0.41) predict the highest rates,

while Africa (0.38) and other emerging regions predict lower probabilities, signalling geographic access barriers. The Category $\times$ Age Group heatmap shows the 21–25 cohort achieves the highest internship predicted rate (0.44), while learners aged 16–20 predict the lowest (0.38).

3.7 Temporal Pipeline Analysis




Figure 5: 2024 Monthly Application Volume with Q4 Outreach Window recommendation (orange shading).

Q1 2024 (January + February) contributed 2,829 applications, approximately 10× the monthly baseline of ~185 applications. This surge demonstrates the platform's capacity to generate dramatically higher application volumes when targeted outreach is conducted. Q3 (September) and Q4 (October–November) currently represent under-utilised periods that, with deliberate campaign activity, can prime the pipeline for a strong Q1 intake in the following year.

4. Actionable Recommendations

The following six recommendations are derived directly from model outputs, probability scores, heatmaps, and temporal analysis. They are ordered by estimated impact and urgency.

Priority	Area	Evidence	Recommendation
Critical	Internship Capacity	63.3% of applications, 21.5% Active rate	Expand placement slots; introduce associate-level feeder internships; open applications only after capacity is confirmed
High	Outreach Timing	Q1 2024 surge = 10× baseline	Launch Q4 (Oct–Nov) campaigns annually; Africa-specific Q2 outreach when competition from NA/Asia is lower
High	Africa & Gender Equity	Africa: 14.8% of apps, lowest internship prob; female Active gap in Africa	Design Africa-specific opportunities; partner with regional women-in-tech organisations; set equity KPI targets
Medium	Drop-Out Prevention	99.4% of 617 drop-outs in Internship	Deploy model probability scores to trigger week 2 and week 4 check-in emails for learners scoring below 40%
Medium	Course Pipeline	Courses: 98%+ predicted Active rate	Make course completion a mandatory pre-qualifier for internship applications to build

Priority	Area	Evidence	Recommendation
			learner readiness and reduce rejection
Low	Model Deployment	3 models trained, GB AUC 91.99%	Deploy Gradient Boosting as live scoring engine; retrain quarterly; extend to multi-class drop-out prediction in Phase 4

4.1 Application Pipeline Improvements

- **Capacity-first intake:** Never open internship applications before confirming the number of available placement slots.
- **Pre-screening dashboard:** Display predicted participation probability to applicants at submission. Those below 50% receive automatic course recommendations.
- **Waitlist conversion:** Re-engage Inactive learners with probability 40–60% when new slots open, before a public round is launched.
- **Manager dashboard:** Provide programme staff with a real-time view of applications scored by predicted probability, segmented by region and category.

4.2 Scalable Solutions for Diverse Learner Profiles

- **21–25 learners:** Prioritise this cohort for internship capacity expansion, highest volume and highest predicted engagement.
- **16–20 learners:** Route to Courses and Events first, with internship eligibility unlocked after completing two or more platform courses.
- **African learners:** Create region-specific opportunities (remote-friendly, Africa-partner aligned); launch Q2 to avoid peak North American competition.

- **African women specifically:** Targeted outreach stream through regional women-in-technology and professional associations.
- **31+ learners:** Design career-change and upskilling internship tracks distinct from undergraduate pathways.

5. Measurement Framework

5.1 Four-Phase Deployment Roadmap

Phase	Timeline	Action	Success Metric
Phase 1	Month 1	Deploy Gradient Boosting model as a scoring engine; tag all applications with predicted Active probability	Scoring system live; 100% of applications tagged
Phase 2	Months 2–3	Use probability scores to trigger automated learner check-in emails; learners <50% probability receive course recommendations	Response rate >20% on check-in emails; course enrolment uplift
Phase 3	Months 4–6	Retrain model quarterly; add prior engagement count and major-opportunity alignment score as new features	CV AUC maintained above 91%; new features show positive importance
Phase 4	Months 7–12	Extend to multi-class prediction (Active / Completed / Dropped Out); enable drop-out risk scoring	Drop-out rate reduced by $\geq 50\%$ vs current 7.2% baseline

5.2 Key Performance Indicators

KPI	Current Baseline	12-Month Target	Measurement Method
Overall Active Participation Rate	47.2%	>60%	Monthly platform dashboard; ratio of Active to total applications

KPI	Current Baseline	12-Month Target	Measurement Method
Internship Rejection Rate	41.7%	<25%	Monthly rejections / total internship applications
Internship Drop-Out Rate	7.2%	<3%	Monthly drop-outs / total started
African Learner Internship Active Rate	~38% predicted	> 50%	Quarterly Active / Applied filtered by Region = Africa
African Female Active Rate	Below male average	Gender parity $\pm 5\%$	Quarterly gender $\times$ region cross-tabulation
Q4 Campaign Volume Uplift	~185 apps/month (baseline)	≥ 600 apps/month in Q4	Monthly application count during Oct–Nov outreach
Model CV AUC (post-retrain)	91.12% (GB)	$\geq 91\%$	Quarterly retraining pipeline; logged CV score

6. Supporting Visualisations

6.1 Model Performance — ROC Curves




Figure 6: ROC Curves — All three models. AUC: Gradient Boosting 0.920, Random Forest 0.917, Logistic Regression 0.909.

6.2 Confusion Matrices




Figure 7: Confusion Matrices — test set ($n=1,663$). Symmetric error distribution confirms no class bias.

6.3 Refined Participation Dashboard




Figure 8: Stacked participation status (left) and Active rate by category (right). Internship inverse demand-quality relationship clearly visible.

6.4 Gender × Region Active Rate




Figure 9: Active Participation Rate by Region and Gender. Africa shows widest gender gap — a key equity finding.

6.5 Cross-Validation AUC




Figure 10: 5-Fold CV AUC-ROC — Mean ± Standard Deviation. All models generalise reliably.

7. Conclusion

Over four weeks, Team 8 transformed 8,558 raw participation records into a deployment-ready machine learning model, 22 opportunity probability scores, and a prioritised programme improvement plan. The work demonstrates that the RIT platform has strong learner demand, but structural capacity constraints, particularly in the Internship category, prevent that demand from translating into proportional engagement.

The Gradient Boosting model (AUC-ROC 91.99%) is ready for deployment as a live application scoring engine. The six recommendations, anchored in expanding internship capacity, launching targeted Q4 outreach, addressing the Africa and gender equity gap, preventing drop-outs, building a course pre-qualification pipeline, and deploying the ML model in four phases, provide a concrete, evidence-based roadmap for the next cohort cycle.

The measurement framework, with specific KPI targets and a quarterly tracking cadence, ensures that the improvements proposed in this report can be monitored, evaluated, and iterated upon by Excelerate programme managers. From data to insight to action, this is the full arc Team 8 has delivered.